

# Balanced Trees (*Red-Black* Trees)

Warren Kim

# Applications

Red-Black Trees have a variety of applications. Some include:

- (i) Database indexing (2-4 trees<sup>1</sup>  $\simeq$  *red-black* trees)[e.g. VoltDB]

---

<sup>1</sup>2-4 (sometimes called 2-3-4) trees are a subset of  $B^+$ -trees.

# Applications

Red-Black Trees have a variety of applications. Some include:

- (i) Database indexing (2-4 trees<sup>1</sup>  $\simeq$  *red-black* trees)[e.g. VoltDB]
- (ii) Linux CPU scheduler (Completely Fair Scheduler)

---

<sup>1</sup>2-4 (sometimes called 2-3-4) trees are a subset of  $B^+$ -trees.

# Applications

Red-Black Trees have a variety of applications. Some include:

- (i) Database indexing (2-4 trees<sup>1</sup>  $\simeq$  *red-black* trees)[e.g. VoltDB]
- (ii) Linux CPU scheduler (Completely Fair Scheduler)
- (iii) Linux Virtual Memory Areas (VMA)

---

<sup>1</sup>2-4 (sometimes called 2-3-4) trees are a subset of  $B^+$ -trees.

# Applications

Red-Black Trees have a variety of applications. Some include:

- (i) Database indexing (2-4 trees<sup>1</sup>  $\simeq$  *red-black* trees)[e.g. VoltDB]
- (ii) Linux CPU scheduler (Completely Fair Scheduler)
- (iii) Linux Virtual Memory Areas (VMA)
- (iv) STL Data Structures (e.g. C++'s `std::map`, Java's `HashMap`)

---

<sup>1</sup>2-4 (sometimes called 2-3-4) trees are a subset of  $B^+$ -trees.

# Applications

Red-Black Trees have a variety of applications. Some include:

- (i) Database indexing (2-4 trees<sup>1</sup>  $\simeq$  *red-black* trees)[e.g. VoltDB]
- (ii) Linux CPU scheduler (Completely Fair Scheduler)
- (iii) Linux Virtual Memory Areas (VMA)
- (iv) STL Data Structures (e.g. C++'s `std::map`, Java's `HashMap`)
- (v) Graph algorithm optimizations (for AI/ML)[e.g. K-mean clustering]

---

<sup>1</sup>2-4 (sometimes called 2-3-4) trees are a subset of  $B^+$ -trees.


# Motivation

- (i) Raw binary search tree performance is highly dependant on input order
- (ii) We want to ensure  $\mathcal{O}(\log n)$  performance


## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$



$\{1, 7, 2, 6, 3, 5, 4\}$







## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




## Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$



$\{1, 7, 2, 6, 3, 5, 4\}$







# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$




$\{1, 7, 2, 6, 3, 5, 4\}$




# Example

Suppose we have the input set  $\{1, 2, 3, 4, 5, 6, 7\}$  and consider two input orders:

$\{1, 2, 3, 4, 5, 6, 7\}$



$\{1, 7, 2, 6, 3, 5, 4\}$



# Intuition

*How do we balance binary trees?*

# Intuition

*How do we balance binary trees?*

We can *dynamically* balance the tree!

# Intuition

*How do we balance binary trees?*

We can *dynamically* balance the tree!

We want to enforce a set of well-defined conditions. We can achieve this by adding additional member variables in our `Node` struct.

# Definition and Properties

## Definition

*A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.*



# Definition and Properties

## Definition

A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.

- (i) *Color*: Every node is either **red** or **black**


```
typedef enum Color { RED, BLACK };  
  
struct Node {  
    Color color;  
    int data;  
    Node *left;  
    Node *right;  
    Node *parent;  
};
```

# Definition and Properties

## Definition

A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.

- (i) *Color*: Every node is either **red** or **black**
- (ii) *External*: All nil nodes are **black**





# Definition and Properties

## Definition

A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.

- (i) *Color*: Every node is either **red** or **black**
- (ii) *External*: All nil nodes are **black**
- (iii) *Internal*: A **red** node does not have a **red** child




# Definition and Properties

## Definition


A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.

- (i) *Color*: Every node is either **red** or **black**
- (ii) *External*: All nil nodes are **black**
- (iii) *Internal*: A **red** node does not have a **red** child
- (iv) *Depth*: Every path from the root to *any* leaf node passes through the same number of **black** nodes
- (v) *Root*: The root node is always **black**




# Depth Property

(iv) *Depth*: Every path from the root to *any* leaf node passes through the same number of **black** nodes




# Depth Property

(iv) *Depth*: Every path from the root to *any* leaf node passes through the same number of **black** nodes



# Depth Property

(iv) *Depth*: Every path from the root to *any* leaf node passes through the same number of **black** nodes



# Definition and Properties

## Definition

A **red-black** tree is a type of **self-balancing** binary search tree that guarantees  $\mathcal{O}(\log n)$  search, insertion, and deletion operations.

- (i) *Color*: Every node is either **red** or **black**
- (ii) *External*: All nil nodes are **black**
- (iii) *Internal*: A **red** node does not have a **red** child
- (iv) *Depth*: Every path from the root to *any* leaf node passes through the same number of **black** nodes
- (v) *Root*: The root node is always **black**



# Insertion

Suppose we have a node  $z$  to insert into our *red-black* tree. Then,

# Insertion

Suppose we have a node  $z$  to insert into our *red-black* tree. Then,

- (i) Like a BST, insert  $z$ .

# Insertion

Suppose we have a node  $z$  to insert into our *red-black* tree. Then,

- (i) Like a BST, insert  $z$ .
- (ii) Color  $z$  *red*.

# Insertion

Suppose we have a node  $z$  to insert into our *red-black* tree. Then,

- (i) Like a BST, insert  $z$ .
- (ii) Color  $z$  *red*.
- (iii) Fix double *red* violations, if any.




# Insertion

Suppose we have a node  $z$  to insert into our *red-black* tree. Then,

- (i) Like a BST, insert  $z$ .
- (ii) Color  $z$  *red*.
- (iii) Fix double *red* violations, if any.
- (iv) Recursively fix violations upward.




## Double Red Violations

Recall *Property (iii)*: A **red** node does not have a **red** child. All of the diagrams shown below are examples of *invalid red-black* trees.



# Double Red Violations




Recall *Property (iii)*: A **red** node does not have a **red** child. All of the diagrams shown below are examples of *invalid red-black* trees.



There are two cases:

# Double Red Violations

Recall *Property (iii)*: A **red** node does not have a **red** child. All of the diagrams shown below are examples of *invalid red-black* trees.




There are two cases:

- (i) **Recolor**: If both the *parent* and *uncle* are **red**, perform a *recolor*.



# Recolor



where


$z$  is the new node

$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

# Recolor



where


$z$  is the new node

$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

# Recolor



where

$z$  is the new node




$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

# Double Red Violations

Recall *Property (iii)*: A **red** node does not have a **red** child. All of the diagrams shown below are examples of *invalid red-black* trees.



There are two cases:


- (i) **Recolor**: If both the *parent* and *uncle* are **red**, perform a *recolor*.
- (ii) **Restructure**: If the *parent* is **red** but the *uncle* is **black**, perform a *tri-node restructure*.

# Tri-Node Restructure

There are four cases:

- (i)* Left-Left
- (ii)* Right-Right
- (iii)* Left-Right
- (iv)* Right-Left

## Case: Left-Left



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Left-Left



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Left-Left



where

$z$  is the new node

$v$  is  $z$ 's parent


$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree



## Case: Right-Right



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Left-Right



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Left-Right



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Left-Right



where

$z$  is the new node


$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree

## Case: Right-Left



where

$z$  is the new node

$v$  is  $z$ 's parent

$u$  is  $z$ 's grandparent

$w$  is  $z$ 's uncle

$T$  is  $v$ 's subtree


# Time and Space Complexities

***Insertion:***  $\mathcal{O}(\log n)$



***Deletion:***  $\mathcal{O}(\log n)$

***Search:***  $\mathcal{O}(\log n)$

## Red-Black Tree: Example




## Red-Black Tree: Example






## Red-Black Tree: Example




## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




## Red-Black Tree: Example





*Case: Recolor*

## Red-Black Tree: Example



*Case: Recolor*


## Red-Black Tree: Example




*Case: Recolor*



## Red-Black Tree: Example




## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example




## Red-Black Tree: Example





*Case: Recolor*

## Red-Black Tree: Example



*Case: Recolor*


## Red-Black Tree: Example



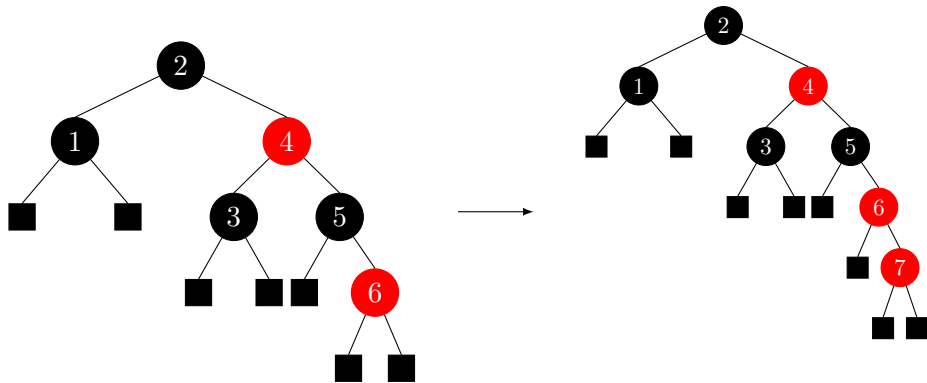
*Case: Recolor*



## Red-Black Tree: Example




## Red-Black Tree: Example




*Case: Right-Right*

## Red-Black Tree: Example



## Red-Black Tree: Example



End

Thank you!

# Appendix

Below are slides that didn't make the cut :(

# Corollaries

## Proposition

*If a node  $n$  has exactly one child,  $c$ , then (a)  $c$  is **red**, (b)  $n$  is **black**, and (c)  $c$  has no children.*

*Proof.* Suppose we have a valid red-black tree. Consider a node  $n$  with exactly one child. Without loss of generality, choose  $n$ 's left node to be the child and call it  $c$ .

- (a)  $n$  passes through no **black** nodes on the right side by assumption. If  $c$  were **black**, then  $n$  would pass through 1 **black** node, a contradiction since this violates the *depth property*.
- (b) By (a),  $n$ 's child is **red** and by the *internal property*,  $n$  is **black**.
- (c) Since  $n$  passes through no **black** nodes on the right side by assumption,  $n$  cannot pass through any **black** nodes on the left side by the *depth property*. Then, since  $c$  is **red** by (a),  $c$  has only nil nodes, which are **black** by the *external property*. □

# Height of a Red-Black Tree

## Theorem

*A red-black tree with  $n$  nodes has a height  $h$  that is  $\mathcal{O}(\log n)$ .*

*Proof.* Suppose we have an arbitrary (valid) red-black tree. Let  $b$  be the number of **black** nodes on the shortest path from root to any leaf. In the worst case, the shortest path alternates between **red** and **black** nodes and thus has a height of  $2b$ . Then,  $h$  is bounded above by  $2b$ ; that is,  $h \leq 2b$ . There are  $2^b - 1 \leq n$  nodes in this tree. Solving for  $b$ , we get  $b \leq \log(n + 1)$ . Substituting  $b$ , we get  $b \leq \log(n + 1) \leq h \leq 2b \leq 2 \log(n + 1)$  so  $h$  is bounded below by  $\log(n + 1)$  and above by  $2 \log(n + 1)$ ; that is,  $\log(n + 1) \leq h \leq 2 \log(n + 1)$ . So,  $h$  is  $\mathcal{O}(\log n)$ .  $\square$



# Deletion

Suppose we have a node  $z$  to delete from our red-black tree.

- (i) If  $z$  is a **red** leaf node, simply remove  $z$ .
- (ii) If  $z$  only has a left child  $s$ , swap *values* and remove  $s$ .
- (iii) If  $z$  has a right child, swap its *value* with its in-order successor  $s$  and remove  $s$ .


If there is a double **black** in (ii) or (iii), perform a *fixup* at  $y$ 's original position.

# Deletion Fixup

There are three cases:

- (i)  $y$  has a **red** child. Then, we perform a *restructure*.
- (ii)  $y$ 's sibling  $w$  is **black**. Then, we perform a *recolor*.
- (iii)  $y$ 's sibling  $w$  is **red**. Then, we perform an *adjustment* followed by either case (ii) or (iii).

# Restructure



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Restructure



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Restructure



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Restructure



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Recolor



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Recolor



where

$z$  is the node to delete


$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child



# Adjustment



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

# Adjustment



where


$z$  is the node to delete

$y$  is  $z$ 's sibling

$x$  is  $y$ 's parent

$c$  is  $y$ 's child

Case:  $y$  is a Leaf Node or has Two ***Black*** Children




where

$y$  is the node to delete

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

Case:  $y$  is a Leaf Node or has Two **Black** Children




where

$y$  is the node to delete

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

Case:  $y$  has a *Red* Child



where



$y$  is the node to delete

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

$z$  is  $y$ 's child

Case:  $y$  has a *Red* Child



where


$y$  is the node to delete

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

$z$  is  $y$ 's child

Case:  $w$  has a *Red* Child



where


$y$  is the node to delete

$z$  is  $y$ 's child

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

Case:  $w$  has a *Red* Child



where

$y$  is the node to delete


$z$  is  $y$ 's child

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent



Case:  $w$  has a *Red* Child



where


$y$  is the node to delete

$z$  is  $y$ 's child

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

Case:  $w$  is *Red*



where

$y$  is the node to delete


$z$  is  $y$ 's child

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

$T$  is  $w$ 's subtree

Case:  $w$  is *Red*



where

$y$  is the node to delete

$z$  is  $y$ 's child

$w$  is  $y$ 's sibling

$x$  is  $y$ 's parent

$T$  is  $w$ 's subtree

## Red-Black Tree Example

